

# CLIP Embeddings as Product Characteristics in BLP Merger Simulation

Evidence from the hello→Colgate Acquisition in the US Toothpaste Market

Burak Uzun

May 2026

---

## 1 Data and Panel Construction

The dataset is an ASIN x quarter panel covering the Amazon US toothpaste market from 2019Q1 to 2022Q3 (19 quarters, 159 ASINs, 1,164 observations). Market shares are constructed using revealed-preference purchase data. Market size  $M_t$  is set to  $(1 + \lambda) \cdot \sum_j q_{jt}$  with  $\lambda = 13.93$ , implying an outside-good share of approximately 93%.

CLIP (Contrastive Language-Image Pre-training; Radford et al. 2021) embeddings are computed for each ASIN using its product image and listing text. A joint 512-dimensional embedding is reduced to five principal components (PC1–PC5), which together explain 71.2% of the within-cluster variance. Five clusters are obtained via  $K$ -means ( $K = 5$ , 50 restarts, seed 42).

---

## 2 Demand Specification

### 2.1 Berry (1994) Linearisation

All three models are estimated via the Berry (1994) nested-logit linearisation:

$$\log \frac{s_{jt}}{s_{0t}} - \rho \log s_{j|g,t} = \alpha p_{jt} + \beta' \mathbf{x}_{jt} + \text{FE}_t + \xi_{jt}$$

where  $s_{j|g,t} = s_{jt} / \sum_{k \in g} s_{kt}$  is the within-nest share,  $\rho \in [0, 1)$  is the nesting intensity, and  $\text{FE}_t$  are quarter fixed effects (absorbed by within-quarter demeaning).

## 2.2 Price Coefficient Calibration

Valid cost-side instruments for price are unavailable in this high-frequency online retail dataset (see Section 5). Following the recommendation of Berry (1994) and standard practice in applied antitrust simulation,  $\alpha$  is calibrated to match a target median own-price elasticity:

$$\alpha = \frac{\varepsilon_{\text{target}}}{\text{median}(p_{jt}(1 - s_{jt}))} = \frac{-2.62}{8.4502} = -0.310$$

The target  $\varepsilon_{\text{target}} = -2.62$  corresponds to the brand-level CPG meta-analysis mean of Bijmolt, van Heerde & Pieters (2005), consistent with structural BLP estimates for analogous packaged-goods categories (Nevo 2001; Draganska & Jain 2006).

**Sequential calibration note.** Since  $\alpha$  is calibrated at  $\rho = 0$ , the target elasticity  $-2.62$  is the logit median. At the profiled nesting parameters ( $\rho^* = 0.35$  for M2,  $\rho^* = 0.67$  for M3), the nested-logit formula amplifies own-price sensitivity by approximately  $1/(1 - \rho^*)$ , yielding implied median elasticities of approximately  $-4.0$  (M2) and  $-7.9$  (M3).

## 2.3 $\rho^*$ Selection — Concentrated OLS Profile

For each  $\rho$  on a fine grid  $\{0, 0.01, \dots, 0.90\}$ , the adjusted dependent variable is formed and  $\beta(\rho)$  is estimated by OLS. The profiled  $\rho^*$  minimises the residual sum of squares:

$$\tilde{y}_{jt}(\rho) = \log \frac{s_{jt}}{s_{0t}} - \alpha p_{jt} - \rho \log s_{j|g,t}, \quad \rho^* = \arg \min_{\rho} \left\| \tilde{y}(\rho) - \mathbf{X} \hat{\beta}(\rho) \right\|^2$$

## 3 Tables

Table 1: Summary Statistics by CLIP Cluster

| Cluster | Dom. Brand     | $N$ ASINs | Price (USD) |             |            | Mkt. Share (x1000)    |                         |                        | $\bar{w}$ | $\tilde{w}$ | $\sigma_w$ |
|---------|----------------|-----------|-------------|-------------|------------|-----------------------|-------------------------|------------------------|-----------|-------------|------------|
|         |                |           | $\bar{p}$   | $\tilde{p}$ | $\sigma_p$ | $\bar{s} \times 10^3$ | $\tilde{s} \times 10^3$ | $\sigma_s \times 10^3$ |           |             |            |
| C1      | Sensodyne      | 29        | 10.26       | 9.63        | 5.58       | 1.068                 | 0.770                   | 0.848                  | 249       | 249         | 145        |
| C2      | Colgate        | 33        | 8.00        | 8.39        | 3.65       | 1.303                 | 0.817                   | 1.259                  | 537       | 399         | 296        |
| C3      | Crest          | 29        | 8.72        | 6.88        | 4.18       | 1.267                 | 0.757                   | 1.282                  | 375       | 431         | 128        |
| C4      | Tom's of Maine | 13        | 9.85        | 9.76        | 2.96       | 1.170                 | 0.770                   | 0.972                  | 374       | 370         | 94         |
| C5      | Other (hello)  | 55        | 9.46        | 7.86        | 6.84       | 0.839                 | 0.670                   | 0.736                  | 171       | 141         | 105        |

*Notes:*

Panel: 1,164 ASIN x quarter observations, 19 quarters (2019Q1–2022Q3). Mean/Median/SD of price in USD, market share scaled x1000, packa

Table 2: Demand Parameter Estimates — Three Models

| Variable                     | M1: Logit ( $\rho=0$ ) | M2: Brand-Nested ( $\rho^*=0.35$ ) | M3: CLIP-Nested ( $\rho^*=0.67$ ) |
|------------------------------|------------------------|------------------------------------|-----------------------------------|
| $\alpha$ (price)             | <b>-0.310 (a)</b>      | <b>-0.310 (a)</b>                  | <b>-0.310 (a)</b>                 |
| Package size (g)             | 0.0037*** (0.0003)     | 0.0037*** (0.0003)                 | 0.0034*** (0.0003)                |
| CLIP $PC_1$                  | -0.4034*** (0.0424)    | -0.2869*** (0.0418)                | -0.3280*** (0.0407)               |
| CLIP $PC_2$                  | 0.2033*** (0.0374)     | 0.2681*** (0.0368)                 | 0.2103*** (0.0359)                |
| CLIP $PC_3$                  | -0.1667*** (0.0386)    | -0.1479*** (0.0380)                | -0.2854*** (0.0370)               |
| CLIP $PC_4$                  | 0.2115*** (0.0399)     | 0.1574*** (0.0393)                 | 0.2707*** (0.0382)                |
| CLIP $PC_5$                  | -0.1652*** (0.0390)    | -0.1243*** (0.0384)                | -0.1745*** (0.0374)               |
| Other brand                  | 0.0849 (0.1327)        | 0.4590*** (0.1307)                 | 0.2641** (0.1272)                 |
| Import x freight shock       | 0.6767*** (0.2342)     | 0.6709*** (0.2306)                 | 0.7461*** (0.2245)                |
| $\rho^*$                     | 0.000 (b)              | 0.350 (b)                          | 0.670 (b)                         |
| Implied median $\varepsilon$ | -2.62                  | $\approx -4.0$                     | $\approx -7.9$                    |
| RSS drop vs $\rho=0$         | —                      | 3.1%                               | 8.1%                              |

Notes:

\*\*\* p&lt;0.01, \*\* p&lt;0.05, \* p&lt;0.10. FEs by quarter. N=1,164. (a) calibrated alpha; (b) grid-searched rho\*. Secs. 2.2-2.3.

Table 3: Merger Counterfactuals: hello-to-Colgate Acquisition (2020Q1)

| Brand                                              | Obs. $p$    | MC          | Nash (pre)  | Nash (post) | $\Delta p$   | $\Delta p$ %  |
|----------------------------------------------------|-------------|-------------|-------------|-------------|--------------|---------------|
| <b>M1: Logit (<math>\rho=0</math>)</b>             |             |             |             |             |              |               |
| hello                                              | <b>6.43</b> | <b>3.20</b> | <b>7.10</b> | <b>7.16</b> | <b>0.064</b> | <b>1.233%</b> |
| Colgate                                            | <b>8.45</b> | <b>5.16</b> | <b>8.29</b> | <b>8.30</b> | <b>0.007</b> | <b>0.104%</b> |
| Tom's of<br>Maine                                  | 10.47       | 7.17        | 10.70       | 10.71       | 0.006        | 0.069%        |
| Crest                                              | 8.87        | 5.60        | 9.18        | 9.18        | 0.000        | 0.000%        |
| Sensodyne                                          | 10.45       | 7.17        | 10.05       | 10.05       | 0.000        | 0.000%        |
| <b>M2: Brand-Nested (<math>\rho^*=0.35</math>)</b> |             |             |             |             |              |               |
| hello                                              | <b>6.43</b> | <b>3.20</b> | <b>7.10</b> | <b>7.17</b> | <b>0.065</b> | <b>1.275%</b> |
| Colgate                                            | <b>8.45</b> | <b>5.16</b> | <b>8.30</b> | <b>8.30</b> | <b>0.006</b> | <b>0.104%</b> |
| Tom's of<br>Maine                                  | 10.47       | 7.17        | 10.71       | 10.71       | 0.006        | 0.069%        |
| Crest                                              | 8.87        | 5.60        | 9.18        | 9.18        | 0.000        | 0.000%        |
| Sensodyne                                          | 10.45       | 7.17        | 10.07       | 10.07       | 0.000        | 0.000%        |
| <b>M3: CLIP-Nested (<math>\rho^*=0.67</math>)</b>  |             |             |             |             |              |               |
| hello                                              | <b>6.43</b> | <b>5.20</b> | <b>7.05</b> | <b>7.07</b> | <b>0.020</b> | <b>0.404%</b> |
| Colgate                                            | <b>8.45</b> | <b>5.88</b> | <b>8.59</b> | <b>8.59</b> | <b>0.002</b> | <b>0.018%</b> |
| Tom's of<br>Maine                                  | 10.47       | 7.19        | 10.72       | 10.72       | 0.002        | 0.021%        |
| Crest                                              | 8.87        | 5.63        | 9.18        | 9.18        | 0.000        | 0.000%        |
| Sensodyne                                          | 10.45       | 7.28        | 10.29       | 10.29       | 0.000        | 0.000%        |

Notes:

All prices USD; post-merger averages (2020Q1–2022Q3). Bold = merger parties. Nash (pre): Bertrand-Nash under pre-merger ownership. Nash

## 4 Figures




Figure 1: Product space from CLIP embeddings. Each point is one ASIN positioned by its first two principal components of the joint image+text embedding. Cluster centroids are labelled C1–C5. hello (red diamonds, C5) and Colgate (dark blue triangles, C2) occupy different clusters, implying a lower diversion ratio and a smaller predicted merger effect under CLIP nesting relative to plain logit.




Figure 2: Merger simulation price trajectories (Model 3, CLIP-Nested,  $\rho^* = 0.67$ ). Top panel: Nash equilibrium prices under merged and pre-merger ownership. Observed prices are omitted from the top panel because their quarter-to-quarter variation (\$4–\$10) would compress the counterfactual divergence. Bottom panel: merger price effect in percent,  $(p^{\text{merger}} - p^{\text{no-merger}})/p^{\text{no-merger}} \times 100$ . hello (red) shows a positive and time-varying effect; Colgate (blue) remains near zero.

## 5 Bertrand–Nash Merger Simulation

The Bertrand–Nash merger simulation follows the standard approach in Berry (1994) and Nevo (2000). Marginal costs are recovered from the pre-merger first-order condition:

$$\mathbf{mc} = \mathbf{p} + (\Omega_{\text{pre}} \circ \Delta)^{-1} \mathbf{s}$$

where  $\Omega[j, k] = \mathbf{1}[\text{firm}_j = \text{firm}_k]$  and  $\Delta$  is the demand Jacobian with elements:

$$\Delta_{jj} = \alpha s_j \left[ \frac{1}{1-\rho} - \frac{\rho}{1-\rho} s_{j|g} - s_j \right], \quad \Delta_{jk}^{\text{same}} = \alpha s_j \left[ -\frac{\rho}{1-\rho} s_{k|g} - s_k \right], \quad \Delta_{jk}^{\text{diff}} = -\alpha s_j s_k$$

Marginal costs are averaged per ASIN across pre-merger quarters and held fixed. Two Nash equilibria are computed via the damped fixed-point iteration of Nevo (2000):

$$\mathbf{p}_{n+1} = \lambda \left[ \mathbf{mc} - (\Omega \circ \Delta(\mathbf{p}_n))^{-1} \mathbf{s}(\mathbf{p}_n) \right] + (1 - \lambda) \mathbf{p}_n$$

with  $\lambda = 0.4$ , convergence tolerance  $10^{-8}$ , and up to 2,000 iterations. The post-merger simulation replaces  $\Omega_{\text{pre}}$  with  $\Omega_{\text{post}}$  (hello ASINs assigned to Colgate’s firm identifier from 2020Q1). The merger price effect is the difference between these two equilibria:  $\Delta p = p^{\text{merger}} - p^{\text{no-merger}}$ .

**Algorithm validation.** The manual R implementation is validated against PyBLP’s `compute_prices()` (Conlon & Gortmaker 2020). For both M1 (plain logit) and M3 (CLIP-nested,  $\rho^* = 0.67$ ), the maximum absolute difference across 847 post-merger observations is below  $2 \times 10^{-8}$ , confirming algorithmic equivalence.

---

## 6 References

- Bijmolt, T.H.A., van Heerde, H.J., & Pieters, R.G.M. (2005). New empirical generalizations on the determinants of price elasticity. *Journal of Marketing Research*, 42(2), 141–156.
- Berry, S.T. (1994). Estimating discrete-choice models of product differentiation. *RAND Journal of Economics*, 25(2), 242–262.
- Conlon, C., & Gortmaker, J. (2020). Best practices for differentiated products demand estimation with PyBLP. *RAND Journal of Economics*, 51(4), 1108–1145.
- Draganska, M., & Jain, D.C. (2006). Consumer preferences and product-line pricing strategies: An empirical analysis. *Marketing Science*, 25(2), 164–174.
- Nevo, A. (2000). A practitioner’s guide to estimation of random-coefficients logit models of demand. *Journal of Economics & Management Strategy*, 9(4), 513–548.
- Nevo, A. (2001). Measuring market power in the ready-to-eat cereal industry. *Econometrica*, 69(2), 307–342.
- Radford, A., Kim, J.W., Hallacy, C., Ramesh, A., Goh, G., Agarwal, S., ... & Sutskever, I. (2021). Learning transferable visual models from natural language supervision. *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 8748–8763.

## 7 Appendix

### 7.1 A: K Sensitivity Analysis

Table 4: K Sensitivity Analysis: CLIP Cluster Count K = 2–10

| $K$      | WCSS         | Silhouette   | $\rho^*$    | RSS Drop    | Neg. MC%    | hello $\Delta p\%$ | Colgate $\Delta p\%$ |
|----------|--------------|--------------|-------------|-------------|-------------|--------------------|----------------------|
| 2        | 1042.0       | 0.199        | 0.60        | 6.1%        | 0.5%        | 1.71%              | 0.015%               |
| 3        | 836.7        | 0.256        | 0.66        | 7.8%        | 0.9%        | 3.49%              | 0.013%               |
| 4        | 641.2        | 0.330        | 0.57        | 5.7%        | 3.0%        | 3.96%              | 0.041%               |
| <b>5</b> | <b>456.5</b> | <b>0.388</b> | <b>0.67</b> | <b>8.1%</b> | <b>1.7%</b> | <b>0.40%</b>       | <b>0.018%</b>        |
| 6        | 365.6        | 0.387        | 0.60        | 6.7%        | 3.4%        | 0.59%              | 0.044%               |
| 7        | 322.0        | 0.368        | 0.48        | 4.6%        | 4.4%        | 0.69%              | 0.055%               |
| 8        | 279.3        | 0.400        | 0.37        | 2.9%        | 4.7%        | 0.82%              | 0.067%               |
| 9        | 247.7        | 0.403        | 0.35        | 2.4%        | 4.7%        | 1.26%              | 0.177%               |
| 10       | 230.7        | 0.373        | 0.40        | 3.8%        | 4.6%        | 0.79%              | 0.065%               |

*Notes:* K-means on 5 CLIP PCs (seed 42, 50 restarts). WCSS = within-cluster sum of squares. Silhouette = average silhouette width (higher = better separation).  $\rho^*$  = argmin RSS on grid  $[0, 0.01, \dots, 0.90]$ . RSS drop =  $100 \times (RSS_0 - RSS^*)/RSS_0$ . Neg. MC = share of ASINs with negative implied marginal cost. hello / Colgate  $\Delta p$  = average post-merger price effect in percent (2020Q1–2022Q3). **Bold/green row = selected specification (K=5).**

### 7.2 B: Model Fit — RMSE and MAE

Table 5: Model Fit: RMSE / MAE vs. Observed Prices (Post-Merger Quarters 2020Q1–2022Q3)

| Brand / Scope | RMSE / MAE (USD)     |                      |                      |
|---------------|----------------------|----------------------|----------------------|
|               | M1: Logit            | M2: Brand-Nested     | M3: CLIP-Nested      |
| <b>All</b>    | <b>1.558 / 0.743</b> | <b>1.556 / 0.746</b> | <b>1.564 / 0.810</b> |
| hello         | 1.579 / 0.836        | 1.580 / 0.837        | 1.533 / 0.799        |
| Colgate       | 1.144 / 0.559        | 1.142 / 0.560        | 1.149 / 0.652        |

*Notes:* Each cell shows RMSE / MAE in USD against observed transaction prices. Nash (merger) prices are used as the model prediction. “All” = all 847 ASIN  $\times$  quarter observations in the post-merger period (2020Q1–2022Q3); hello = 35 obs.; Colgate = 172 obs. Pre-merger marginal costs are held fixed in both Nash simulations.